

Viral Reels Pipeline: Roadmap

The goal isn't just "a video" — it's a **viral, addicting reel that makes money**: a story people can't swipe away from, built efficiently and repeatably. Everything below is organized around that, not just "what the code does."


Three tiers: what's built and working, what's a good next addition, and open questions worth deciding on purpose rather than by default.

- Exists today
- 🔄 Proposed / not built
- 🤖 Worth deciding — not a feature, a choice
- Human review gate

01 HOW A VIDEO ACTUALLY GETS MADE

The real state machine (`pipeline/core/state_machine.py`), not a simplified happy path: every gate can send the job backward, one fork skips a whole leg, and any running stage can fail and be retried right where it left off.

— forward event — conditional fork (skip Sourcing) - - - reject + feedback (loop back)
 - - - manual jump back a gate - - - fail / retry - - - cancel



02 EFFICIENT & STREAMLINED – YOU ASKED FOR THIS EXPLICITLY

"Viral" and "efficient" pull in different directions if you're not careful: more human review = safer and better-matched footage, but slower per video. Here's where the pipeline already scales, and where the real bottleneck sits.

EXISTS

PROPOSED

WORTH DECIDING

Concurrent jobs

Orchestrator locks per-job and runs each in its own async task — two stories can already run side by side.

One-click "next batch"

Pull more assets at Gate 2 without leaving the page or retyping search terms.

Reusable keyword files

A generated keyword file can be reused across retries without regenerating.

Jobs dashboard

Every job's state/age/what's-pending at a glance instead of one terminal per story.

Batch mode

A whole content calendar in, one job per row, all waiting at Gate 1.

Render queue + concurrency cap

MoneyPrinterTurbo renders in one container; two renders at once will currently just slow each other.

Where's the real bottleneck?

It's the three human gates, not compute. "Efficient" may mean fewer decisions per video, not more parallel jobs.

Auto-approve low-risk items?

Skipping review for high-relevance, low-risk assets would speed things up but trades away the safety net that's been built so carefully.

03 HOOK & STORY STRUCTURE

The first 1-3 seconds decide if they swipe.

EXISTS

"Open with a hook"

One line in the script-writer prompt. No formula, no variants, nothing scoring it.

Chronological, fact-grounded script

Written straight from the source text in order — accurate, but not structured for suspense.

PROPOSED

Cold-open rewrite

Lead on the most shocking beat, then rewind — instead of telling it start to finish.

Multiple hook variants

Generate 3-5 hook options per video instead of committing to the model's first attempt.

Shot-planning pass

An LLM reads the finished script and plans what each scene needs, instead of round-robin keywords.

WORTH DECIDING

Accuracy vs. drama

Restructuring for suspense can pull against staying strictly source-grounded — worth a stated line you won't cross, especially for true crime.

Test hooks, don't just pick one

A/B'ing a couple of hook variants beats trusting a single LLM guess at what's "strong."

04 PACING & VISUAL CRAFT

EXISTS MOST OF THIS SHIPPED TODAY

Word-by-word captions

Animated "pop_spring" reveal, not a static sentence — MoneyPrinterTurbo's own feature, just turned on.

PROPOSED

LLM-based clip matching

Replace bag-of-words scene↔photo matching with real judgment — the single biggest fix for "photo doesn't match the sentence."

WORTH DECIDING

Transitions: on or off?

Supported but left off — nobody's watched a rendered example with one yet to judge if it looks good or cheap.

Background music

On by default, kept low under the narration.

3-second cuts

Down from 5 — more cuts per minute, harder to scroll past.

Automatic Ken Burns zoom

Every still photo already gets a slow zoom — no configuration needed, it was already happening.

Cover-frame choice

Pick the swipe-stopping first frame on purpose instead of whatever the render opens on.

Is sequential concat enough?

Whether "whichever clip matched" can ever look as tight as a real edit, or if scenes need more distinct footage options per beat.

05 SOURCING & STORY-AWARE SEARCH

EXISTS

11 source adapters

Free/license-safe by default; every request and kept file is auditable.

Internet Archive fix

No longer silently drops unlicensed-but-real items (this week's fix).

PROPOSED

Openverse aggregator

Flickr Commons + museums, one API, cleaner license metadata.

Newspaper archive

Period headlines/clippings — coverage into the right decade still needs checking.

Generic era B-roll fallback

A deliberate fallback tier so no scene ever ends up with literally nothing.

Vision-based vetting

Today's flags read text/metadata only, never the actual pixels.

WORTH DECIDING

Is free-only capping quality?

Staying 100% free/license-safe may limit image density and variety enough that a small paid-stock or AI-image budget pays for itself once there's real revenue.

06 DISTRIBUTION & CTA

EXISTS NEW TODAY

Title / caption / hashtags

Generated per platform (Reels, TikTok, Shorts) from the final script, written to SOCIAL_POST.md.

PROPOSED

Posting scheduler

Consistent cadence matters for growth on every platform — nothing here manages that yet.

WORTH DECIDING

Per-platform tuning

Beyond the generic template — what actually works differs by platform and changes over time.

Don't trust the copy blindly

The generated title/caption/hashtags reflect general best practice, not real-time knowledge of what's working on these platforms right now — spot-check a few against accounts that are actually performing well before relying on it.

07 FEEDBACK & LEARNING LOOP

EXISTS

Nothing yet

Every choice today — topic, hook, pacing — is a one-shot guess with no data closing the loop.

PROPOSED

Performance ingestion

Pull real watch-time/retention/follow data after posting, feed it back into future choices.

WORTH DECIDING

Full API vs. manual MVP

Platform analytics APIs are often gated behind approval or a follower minimum. A lightweight "paste in the numbers after each post" version may be the realistic first step, not full API integration.

This is probably the highest-leverage gap

Everything else here is a guess about what makes a video good. This is the only thing that would actually tell you.

08 COMPLIANCE & EDITORIAL RISK

EXISTS

Sensitive-content flags

Vetting already flags gore/violence/child-related keywords as high severity for a human to see before approving.

PROPOSED

Nothing built here beyond flags

No platform-policy checker, no disclosure-labeling step.

WORTH DECIDING — VERIFY BEFORE SCALING UP

Demonetization/shadowban risk

True-crime content specifically can trip platform policy depending on how graphic it gets — worth checking each platform's current monetization policy directly rather than assuming, since this changes.

AI-content disclosure

Several platforms now require labeling AI-assisted content — worth verifying current requirements before posting at volume, not something to guess at.

Real people, real families

A legal/license check doesn't cover editorial sensitivity toward victims and families — that stays a human judgment call, every time.

09 SERIES & BRAND IDENTITY

EXISTS

Nothing

No field anywhere tracks a content-calendar "Series" (e.g. "The One Who Survived").

PROPOSED

Series tracking field

On the job model, in the output filename, in the credits.

Consistent intro/outro template

A recognizable open/close format across every video in a series.

WORTH DECIDING

This may out-rank hook rewrites

A consistent, recognizable format is one of the more reliable levers for return viewers and swipe-resistance — worth prioritizing alongside hook work, not after it.